Supplementary Material: Benchmark Definitions, Execution Protocol, and Uncertainty

The Grid-GDP functions are adapted from the Quantum Hamiltonian Descent collection; each instance is defined by its implemented expression and domain rather than its conventional name (SECQUOIA, 2026; Wu Group, n.d.).

S1 Coordinate domains and normalization

For each row of Table S1, both endpoints are included and the physical coordinates are

$$a_i = \ell_x + \frac{i}{30}(u_x - \ell_x), \qquad b_j = \ell_y + \frac{j}{30}(u_y - \ell_y), \qquad i, j \in \{0, \ldots, 30\}. \tag{S1}$$

For the raw functions in Section S2, the minimized objective is

$$F(x, y) = \frac{f(x, y) - L}{U - L}, \qquad L = \min_{i,j} f(a_i, b_j), \quad U = \max_{i,j} f(a_i, b_j). \tag{S2}$$

Thus L and U are obtained by exhaustive preprocessing of the 961 lattice points, not from the continuous optimum. All trigonometric arguments are in radians. Full-precision normalization constants and target coordinates are archived in `grid_inventory.json` (Section S6).

Table S1: Grid-GDP benchmark domains and target-set sizes. Each problem uses a 31×31 lattice with 961 configurations and 62 Boolean indicators. n_{target} counts lattice points with normalized objective $F \leq 10^{-5}$.

Function	x domain	y domain	n_{target}
Ackley	$[-5, 5]$	$[-5, 5]$	1
Alpine 1	$[0, 10]$	$[0, 10]$	1
Alpine 2	$[0, 1000]$	$[0, 1000]$	1
Bohachevsky 2	$[-100, 100]$	$[-100, 100]$	1
Three-hump camel	$[-5, 5]$	$[-5, 5]$	1
Csendes (squared variant)	$[-1, 1]$	$[-1, 1]$	36
Deflected corrugated spring	$[0, 10]$	$[0, 10]$	1
Drop-Wave	$[-5.12, 5.12]$	$[-5.12, 5.12]$	1
Easom	$[-10, 10]$	$[-10, 10]$	1
Griewank	$[-600, 600]$	$[-600, 600]$	1
Hölder table	$[-10, 10]$	$[-10, 10]$	4
Hosaki	$[0, 5]$	$[0, 6]$	1
Lévy	$[-10, 10]$	$[-10, 10]$	2
Lévy 13	$[-10, 10]$	$[-10, 10]$	4
Michalewicz	$[0, \pi]$	$[0, \pi]$	1
Rastrigin	$[-5.12, 5.12]$	$[-5.12, 5.12]$	1
Rosenbrock	$[-2, 2]$	$[-1, 3]$	2
Shubert (historical variant)	$[0, 10]$	$[0, 10]$	1
Styblinski–Tang	$[-5, 5]$	$[-5, 5]$	1
Sum of squares	$[-10, 10]$	$[-10, 10]$	1
Xin-She Yang 3 (historical variant)	$[-20, 20]$	$[-20, 20]$	870

S2 Grid-GDP objective functions

The implemented functions below use the domains in Table S1 (SECQUOIA, 2026).

1. Ackley

$$f(x,y) = -20\exp\left(-0.2\sqrt{\frac{x^2+y^2}{2}}\right) - \exp\left(\frac{\cos(2\pi x)+\cos(2\pi y)}{2}\right) + 20 + e. \tag{S3}$$

2. Alpine 1

$$f(x,y) = |x\sin x + 0.1x| + |y\sin y + 0.1y|. \tag{S4}$$

3. Alpine 2

$$f(x,y) = -\sqrt{xy}\,\sin x \sin y. \tag{S5}$$

4. Bohachevsky 2

$$f(x,y) = x^2 + 2y^2 - 0.3\cos(3\pi x)\cos(4\pi y) + 0.3. \tag{S6}$$

5. Three-hump camel

$$f(x,y) = 2x^2 - 1.05x^4 + \frac{x^6}{6} + xy + y^2. \tag{S7}$$

6. Csendes (squared variant)

$$f(x,y) = c(x) + c(y), \qquad c(t) = \begin{cases} t^6\left[2+\sin(1/t)\right]^2, & t \neq 0, \\ 0, & t = 0. \end{cases} \tag{S8}$$

The GDP uses branch-fixed auxiliaries $r_t = 1/t$ for $t \neq 0$ and $r_t = 0$ otherwise, avoiding division by zero without an epsilon shift.

7. Deflected corrugated spring

With $r = \sqrt{(x-5)^2 + (y-5)^2}$,

$$f(x,y) = 0.1r^2 - \cos(5r). \tag{S9}$$

8. Drop-Wave

$$f(x,y) = -\frac{1+\cos\left(12\sqrt{x^2+y^2}\right)}{0.5(x^2+y^2)+2}. \tag{S10}$$

9. Easom

$$f(x,y) = -\cos x \cos y\, \exp\left(-[(x-\pi)^2 + (y-\pi)^2]\right). \tag{S11}$$

10. Griewank

$$f(x,y) = \frac{x^2+y^2}{4000} - \cos x \cos\left(\frac{y}{\sqrt{2}}\right) + 1. \tag{S12}$$

11. Hölder table

$$f(x,y) = -\left|\sin x \cos y\, \exp\left(\left|1 - \frac{\sqrt{x^2+y^2}}{\pi}\right|\right)\right|. \tag{S13}$$

12. Hosaki

$$f(x,y) = \left(1 - 8x + 7x^2 - \frac{7}{3}x^3 + \frac{1}{4}x^4\right) y^2 \exp(-y). \tag{S14}$$

13. Lévy

Define $w_x = 1 + (x-1)/4$ and $w_y = 1 + (y-1)/4$. Then

$$\begin{aligned} f(x,y) = {} & \sin^2(\pi w_x) \\ & + (w_x - 1)^2\big[1 + 10\sin^2(\pi w_x + 1)\big] \\ & + (w_y - 1)^2\big[1 + \sin^2(2\pi w_y)\big]. \end{aligned} \tag{S15}$$

14. Lévy 13

$$\begin{aligned} f(x,y) = {} & \sin^2(3\pi x) + (x-1)^2\big[1 + \sin^2(3\pi y)\big] \\ & + (y-1)^2\big[1 + \sin^2(2\pi y)\big]. \end{aligned} \tag{S16}$$

15. Michalewicz

$$f(x,y) = -\sin x\, \sin^{20}\left(\frac{x^2}{\pi}\right) - \sin y\, \sin^{20}\left(\frac{2y^2}{\pi}\right). \tag{S17}$$

16. Rastrigin

$$f(x,y) = x^2 - 10\cos(2\pi x) + y^2 - 10\cos(2\pi y) + 20. \tag{S18}$$

17. Rosenbrock

$$f(x,y) = (1-x)^2 + 100(y - x^2)^2. \tag{S19}$$

18. Shubert (historical variant)

$$\begin{aligned} f(x,y) = {} & \big[\cos(2x+1) + 2\cos(3x+2) + 3\cos(4x+3)\big] \\ & \times \big[\cos(2y+1) + \cos(y+2)\big]. \end{aligned} \tag{S20}$$

19. Styblinski–Tang

$$f(x,y) = \frac{1}{78}\,\frac{1}{2}\left(x^4 - 16x^2 + 5x + y^4 - 16y^2 + 5y\right). \tag{S21}$$

20. Sum of squares

$$f(x,y) = x^2 + 2y^2. \tag{S22}$$

21. Xin-She Yang 3 (historical variant)

$$f(x,y) = \exp\left(-\left[\left(\frac{x}{15}\right)^6 + y^6\right]\right) + 2\exp\big(-(x^2+y^2)\big)\cos^2 x \cos^2 y. \tag{S23}$$

S3 Executable protocol and implementation settings

Python and NumPy use the job seed, and `RBFOpt` also receives `rand_seed`. `BAM` exposes no separate native random-seed option: its process wrapper seeds Python/NumPy, while the Grid seed identifies the common initial design. `FMQA-SA` uses three NumPy `SeedSequence(seed)` streams for FM initialization, annealing, and uniform unseen-point fallback; each training or annealing invocation draws a seed from its stream. `PyTorch` uses deterministic CPU operations and one thread.

Grid `D-SDA` enumerates L_∞ neighbors lexicographically from $(-1, 0, 1)$, excluding the zero direction and out-of-bounds points; ties select the first point.

The `RBFOpt` options are `rbf=auto`, `global_search_method=genetic`, `do_init_strategy=False`, `ga_base_population_size=400`, and `ga_num_generations=20`. `RBFOpt` 4.3.0 uses an effective population of $400 + \lfloor 20n/5 \rfloor$, where n is the dimension passed to the genetic-search routine, giving 408 and 412 for two and three search variables. `BAM` receives `x0=None` and `max_evals`=B-3; its local models use `ALAMO`/`BARON`, whose internal versions are not reported by the package.

The `FMQA-SA` time limit is checked between operations and through training/sampling callbacks. Each FM uses identity activation, zero bias, and normal initial weights with mean zero and standard deviation 0.01. Training uses Adam (Kingma and Ba, 2015), 1,000 full-batch epochs, mean squared error, and learning rate 0.01. Coordinate ξ_j uses $\lfloor \log_2 u_j \rfloor + 1$ unsigned bits, where $u_j > 0$ is its upper bound; QUBO coefficients are divided by the largest absolute entry of the quadratic matrix. D-Wave's `SimulatedAnnealingSampler` in `dwave-samplers` 1.6.0 uses 100 reads, 1,000 sweeps, and a geometric inverse-temperature schedule (D-Wave Systems, n.d.). All process workers and the Grid `FMQA-SA` worker have a 2,100 s watchdog. The Grid `RBFOpt`/`BAM` trace watchdogs are disabled (`max_run_seconds=0`); archived Grid `D-SDA` uses budget, target, or local-search termination without an additional method-clock limit.

S3.1 Versions and independently timed process evaluations

The recorded Python version is 3.11.16; other versions are given in the main paper (Section 3.4). The recorded Grid `GDPOpt` runs make no NLP solver calls because preprocessing fixes all original variables. The algebraic process models retain their archived initializations, including constructor defaults for `CSTR30`; the updated DFO evaluation table is identified separately in Section S6.

The acceptance rule in `gdp_models/evaluation.py` requires all of the following: a finite objective; termination in {`optimal`, `locallyOptimal`, `globallyOptimal`}; solver status other than `error` or `aborted`; interface status `SOLVESTAT`=1 and `MODELSTAT`$\in \{1, 2\}$; and a finite maximum primal violation no greater than 10^{-5}. Time-limit, iteration-limit, evaluation-limit, and error outcomes are not accepted, even if a finite feasible candidate is available.

The primal check uses unscaled violations of original constraints and initialized-variable bounds, including constraint rows deactivated during preprocessing. Nonfinite objectives, constraint bodies, or initialized-variable values preclude acceptance; unused uninitialized variables are skipped. Preprocessing rejection or worker failure supplies no accepted objective.

The `CSTR` initializer `cstr_box_ipopt.py:initialize_feasible` runs after preprocessing and checks the mass-balance/zero-recycle start at 10^{-8} without changing constraints or bounds. All process evaluations use cold starts; NLP solutions are not transferred between configurations. Objectives, acceptance status, and NLP point costs are linked to the same recorded attempt by model, `CSTR` size, and coordinate, without successful-retry selection or cross-size time substitution.

The CSV is loaded before timing starts; callback logging remains in search time rather than the subtracted lookup time.

S3.2 Observed proposal and initialization behavior

Table S2 pools `FMQA-SA` proposals across budgets and seeds, excluding the shared initialization and including terminal callbacks. Fallback supplies 52.89% of process proposals; the uniform comparison evaluates the complete procedure, without isolating the contributions of learning, annealing, and fallback.

Table S2: Recorded `FMQA-SA` uniform fallback by benchmark family.

Family	Runs	Proposals	Fallbacks	Rate
Grid	1,050	51,450	9,522	18.51%
`CSTR`	60	1,350	369	27.33%
`Column`	10	100	86	86.00%
`Dynamic`	60	1,460	1,063	72.81%
`Small Batch`	10	70	58	82.86%
Process total	140	2,980	1,576	52.89%

S4 Process model scale, units, and discrete decisions

Table S3 counts the original GDP before reformulation or presolve, including all active disjunct branches and fixed continuous variables. `Dynamic` is discretized with 30 finite elements per stage and three Radau points. Binary counts refer to the indicator variables already constructed by the model, before Boolean-to-binary conversion. All nine models have zero additional nonbinary integer variables.

Table S3: Original process-model sizes. n_c: continuous variables; n_b: binary variables; m_a: algebraic constraints; m_l: logical constraints; n_d: disjunctions.

Model	n_c	n_b	m_a	m_l	n_d
`CSTR5`	56	20	100	12	10
`CSTR10`	106	40	195	22	20
`CSTR15`	156	60	290	32	30
`CSTR20`	206	80	385	42	40
`CSTR25`	256	100	480	52	50
`CSTR30`	306	120	575	62	60
`Column`	446	30	637	16	15
`Small Batch`	19	18	34	3	9
`Dynamic`	1647	27	3276	31	9

`CSTR` volume is measured in litres, and `Small Batch` investment uses the source model's dollars without currency-year conversion. The `Column` objective is $10^3(Q_c + Q_b) + 10^3 n_{\text{active}}$, the source model's weighted cost proxy rather than a newly estimated annualized cost. The `Dynamic` objective is $-\int_0^9 x(t)^2\,dt$, with no physical-unit calibration for state or time.

S5 Success-rate uncertainty and completion costs

For each problem, method, and budget, we report target hits h out of ten seeds with pointwise two-sided 95% Clopper–Pearson intervals $[p_L, p_U]$, counting execution failures as non-successes (Clopper and Pearson, 1934). Coverage assumes independent Bernoulli outcomes under the fixed protocol; it is not simultaneous across comparisons. Ten successes give $[0.692, 1]$, and zero give $[0, 0.308]$, so neither observed endpoint establishes a population probability of one or zero. Table S4 reports completion costs for all 14 process conditions.

Per-function Grid results at $B = 100$ and complete seed-level outcomes are available in the analysis files linked in Section S6.

Table S4: Median reconstructed full-run completion costs for all ten paired seeds per condition, including target misses, in seconds.

Case	B	D-SDA	RBFOpt	BAM	FMQA-SA	Uniform
Column	13	35.44	54.60	27.50	41.82	21.69
CSTR5	8	10.67	10.72	10.67	13.92	9.44
CSTR10	15	18.95	21.79	18.54	27.07	17.75
CSTR15	22	29.93	37.59	27.67	42.82	27.84
CSTR20	29	40.68	45.03	39.37	59.10	40.12
CSTR25	36	51.69	62.29	50.47	74.70	51.17
CSTR30	43	53.12	101.96	115.41	101.19	65.20
Small Batch	10	3.84	5.12	4.23	9.09	3.49
Dynamic	12	12.23	12.29	9.95	14.47	7.80
Dynamic	15	15.85	16.10	11.53	18.39	8.71
Dynamic	20	21.30	20.54	14.07	25.23	12.27
Dynamic	30	26.62	34.19	21.91	36.24	19.38
Dynamic	37	26.62	42.42	24.99	46.05	24.38
Dynamic	50	26.62	65.14	33.98	61.31	33.12

These are full-run estimates, not first-hit times.

S6 Data sources and reproducibility

The publication snapshot contains the source code, experimental data, analysis results, and reproduction instructions for the DFO and algebraic comparisons. The Grid `GDPOpt` summary combines original non-error runs with recovery runs for the 46 execution errors, so its counts do not establish a method ranking. The three presolve-disabled LOA diagnostic runs are excluded.

S7 Feasibility and bound consistency of the algebraic results

Of the 45 algebraic runs reported in the main paper (Table 5), 19 have flagged bounds (Table S5) relative to the checked solutions in Table S6. Unrounded values, residuals, and source records are linked in Section S6.

For each minimization model, let f_w be the recorded objective of its best independently checked saved feasible solution, and let L denote a reported lower bound. We apply the same rule to all five methods:

$$\tau_w = 10^{-5}\max\{1, |f_w|\}, \qquad \Delta = L - f_w, \qquad \rho = \Delta/\tau_w, \qquad \dagger \iff \rho > 1. \tag{S24}$$

Missing lower bounds remain missing, not zero.

Saved vectors are loaded into the original GDP to reevaluate objectives and check active constraints, variable bounds, integrality, disjunctions, and logical validity without optimization. Violations are unscaled original-equation residuals, not reformulation residuals. All 43 saved feasible vectors satisfy the 10^{-5} threshold; the largest violation is 9.99×10^{-7}. The two Big-M runs without accepted feasible solutions remain in the 45-run denominator. This numerical check establishes neither exact feasibility nor global optimality.

Twelve of the 19 discrepancies are only 1.09–4.30 times the tolerance and are sensitive to numerical feasibility conventions. These conflicting bounds and `optimal` reports provide no independent global certificate or intrinsic method ranking; their cause in GDPopt or its subsolvers remains unresolved.

Table S5: Flagged bound comparisons, using unrounded $\rho > 1$. C identifies the comparison solution in Table S6; displayed values are rounded. Flags indicate numerical inconsistencies of different magnitudes.

ID	Model	Method	Reported L	$v_{\max}$	C	ρ
S01	`CSTR5`	RIC	3.06201	9.31e-7	C01	1.13†
S02	`CSTR5`	GLOA	3.13019	9.99e-7	C01	2.23e+3†
S03	`CSTR5`	LOA	3.06201	9.31e-7	C01	1.13†
S05	`CSTR5`	Hull	3.06201	9.34e-7	C01	1.09†
S06	`CSTR10`	RIC	2.88954	7.87e-11	C02	4.30†
S07	`CSTR10`	GLOA	3.44794	9.16e-7	C02	1.93e+4†
S08	`CSTR10`	LOA	2.88954	7.87e-11	C02	4.30†
S10	`CSTR10`	Hull	2.88951	9.37e-7	C02	3.54†
S11	`CSTR15`	RIC	2.82939	1.39e-9	C03	1.30†
S12	`CSTR15`	GLOA	3.56190	9.78e-7	C03	2.59e+4†
S13	`CSTR15`	LOA	2.82939	1.39e-9	C03	1.30†
S16	`CSTR20`	RIC	2.79941	4.47e-9	C04	1.67†
S17	`CSTR20`	GLOA	3.61603	8.93e-7	C04	2.92e+4†
S18	`CSTR20`	LOA	2.79941	4.47e-9	C04	1.67†
S21	`CSTR25`	RIC	2.78148	9.72e-11	C05	1.72†
S22	`CSTR25`	GLOA	3.64775	9.83e-7	C05	3.11e+4†
S23	`CSTR25`	LOA	2.78148	9.72e-11	C05	1.72†
S27	`CSTR30`	GLOA	3.66862	9.40e-7	C06	3.25e+4†
S45	`Dynamic`	Hull	-196.931	1.00e-7	C09	7.11e+3†

Additional tightened-feasibility check for `CSTR15`. A Dantzig `IPOPT` 3.14.19 solve retained the archived LOA witness W03's discrete configuration and starting values and targeted an original-GDP

Table S6: Checked saved comparison solutions; displayed values are rounded.

ID	Model	Method	Recorded f_w	$v_{\max}$
C01	`CSTR5`	Big-M	3.06198	9.40e-7
C02	`CSTR10`	Big-M	2.88941	9.76e-7
C03	`CSTR15`	Hull	2.82936	9.47e-7
C04	`CSTR20`	Hull	2.79936	9.90e-7
C05	`CSTR25`	Hull	2.78143	9.90e-7
C06	`CSTR30`	Big-M	2.76957	9.68e-9
C07	`Column`	GLOA	22271.5	1.00e-8
C08	`Small Batch`	GLOA	167428	9.00e-7
C09	`Dynamic`	Big-M	-212.009	1.00e-7

residual of 10^{-8} (SECQUOIA, 2026). W03 is distinct from the current `SCIP` comparison solution C03. Only exactly satisfied constant rows were omitted from the computational clone; the independent check retained all original constraints. The returned objective was 2.82939, with maximum unscaled violation 1.11×10^{-16}. The current `SCIP` GLOA lower bound exceeds this objective by 0.732507, or approximately 2.59×10^4 times the audit tolerance. The discrepancy persists under the tighter check, but its cause remains unresolved.

Additional formulation check for `Dynamic` Hull. The saved Big-M point has objective -212.009, approximately 15.0776 below the Hull lower bound -196.931. It satisfies both the original GDP and a reconstructed Hull formulation at 10^{-5}, with maximum violation approximately 10^{-7} after assigning the auxiliary variables from their definitions. The reconstructed export matches the archived preflight export, but the compiled model from the actual solve was not retained. The check therefore confirms the numerical inconsistency without identifying its solver or interface cause; no optimization was rerun.

S8 Additional candidate-search and sampling comparisons

S8.1 Paired differences and sampling controls

Paired success differences retain the common case, budget and seed. For Grid, the descriptive percentile bootstrap resamples the 21 intact function vectors, each containing ten paired seed differences, rather than treating 210 observations as exchangeable Bernoulli draws. For Process, paired seeds are resampled within each of the 14 conditions; the six related `Dynamic` budgets are not pooled for inference. The seed is 20261020 with 2,000 bootstrap replications and deterministic case/seed ordering. Intervals are exploratory and unadjusted for multiple comparisons; ten-seed boundary cases may yield degenerate intervals, and zero inclusion does not establish equivalence.

Table S7: Selected paired success differences, in percentage points. The first method minus the second defines the sign; Process intervals concern the stated condition only.

Grid

B	Contrast	Difference	95% interval
10	`RBFOpt` – Uniform	13.33	[5.71, 24.29]
100	`RBFOpt` – Uniform	60.00	[42.86, 75.71]
10	`BAM` – `BAM`-prefix control	0.00	[0.00, 0.00]
100	`BAM` – `BAM`-prefix control	10.95	[0.48, 23.81]
10	`FMQA-SA` – Uniform	0.95	[-0.95, 2.86]
100	`FMQA-SA` – Uniform	40.00	[28.57, 52.38]
10	`D-SDA` – `RBFOpt`	-13.81	[-24.29, -6.19]
100	`D-SDA` – `RBFOpt`	-33.33	[-48.10, -19.52]

Process: ten paired seeds per condition

Case	B	Contrast	Outcomes	Difference	95% interval
`Dynamic`	12	`RBFOpt` – `D-SDA`	4/6/0/0	60	[30, 90]
`Small Batch`	10	`FMQA-SA` – Uniform	2/4/2/2	20	[-30, 60]

Outcomes list both successful / first only / second only / neither successful.

S8.2 Grid suite exclusions

Excluding Xin-She Yang 3 and Bohachevsky 2 preserves the main descriptive pattern. For the sampled Bohachevsky 2 lattice, $\cos(3\pi a_i) = 1$ at every coordinate, leaving an objective nearly redundant with sampled Sum of Squares after normalization. Without both functions, $B = 100$ rates are 79.5% for `RBFOpt`, 66.8% for `BAM`, 57.9% for `FMQA-SA`, 54.7% for the `BAM`-prefix control, 42.6% for `D-SDA`, and 17.9% for uniform. At $B = 10$, `BAM` and its prefix control each retain 47.9%, compared with `RBFOpt` 15.8%.

S8.3 Limited reconstructed FMQA diagnostics

Seven FM models were reconstructed locally from archived observations in three selected runs because the original fitted weights and QUBOs were not saved (SECQUOIA, 2026). All 3,328 enumerated binary states passed the recorded energy-conversion checks. Five selected SA proposals minimized their reconstructed models over eligible candidates; two `Dynamic` fallback iterations had no eligible stored SA sample. These selected checks do not verify the historical weights or isolate the contributions of learning, annealing, and fallback.